

Project Name:



Enterprise RAG and Knowledge Intelligence Platform



Architecture and BackEnd Implementation along with Validation done for each Phase

Understood the business requirement and written the prompt to generate the architecture.md file


Pasted the prompt in claude and generating the architecture.md

It had asked 20 questions and I have provided the response based on that architecture file had been created







Architecture.md file had been created




Downloaded .md file is available







Opened the RAG Project Folder in VS Code



Refer the Enterprise-AI-Knowledge-Platform-Blueprint.md architecture file and implement it phase by phase .Each phase should be implemented after my approval only .Once one phase is implemented give me steps to validate and test the same.After tested,I will give approval to implement the next phase [STRICT] IMPLEMENT PHASE BY PHASE AND FOLLOWED MY APPROVAL GO FOR NEXT PHASE



Started Implementation “BackEnd” at first – Implementating Phase 1




Created the dB details

The screenshot displays the MongoDB Cloud Portal interface. The top navigation bar includes tabs for 'Business Requirement Prompt', 'Testleaf', 'Untitled document - Google D...', and 'RAG Data | Cloud: MongoDB'. The main content area shows the 'TestDB' details, including a 'Data Explorer' sidebar and a central pane with a message: 'This collection has no data. It only takes a few seconds to import data from a JSON or CSV file.' The bottom status bar shows 'Alert System Status: All Good' and '©2026 MongoDB, Inc.'.

Implementation of Phase 1 — Foundation is completed

The screenshot displays the Visual Studio Code interface. The top navigation bar includes tabs for 'File', 'Edit', 'Selection', 'View', 'Go', 'Run', and 'RAGPROJECT'. The main content area shows the 'Implement enterprise AI knowledge platform...' project, including a sidebar with 'EXPLORER' and 'OUTLINE' sections. The central pane displays the 'Implement enterprise AI knowledge platform...' project with a message: 'Phase 1 — Foundation: Complete'. The bottom status bar shows 'Chat quota reached'.




Validate the Phase 1



Put your real MongoDB Atlas URI into `backend/.env` (`MONGODB_URI=...`)

1. Follow the Setup and Manual validation sections in `backend/README.md` — 9 concrete curl-based checks against your own Atlas cluster, covering everything above end-to-end


Once you've run through that and you're satisfied, give the go-ahead and I'll start Phase 2 — Ingestion & Chunking.

Doing Setup flow for Phase 1





Run the Phase 1 Validation steps at once and all got the passed



Phase 1 — Foundation Validation... - Run results

Ran today at 02:20:14 PM [View all runs](#)

Source	Environment	Iterations	Duration	All tests	Errors	Avg. Resp. Time
Runner	none	1	9s 40ms	42	0	238 ms

All 42 Passed 42 Failed 0 Skipped 0 Errors 0 Console log

Iteration 1

GET Step 1 — Health Check
http://127.0.0.1:8000/health
200 - 100 ms - 175 B 3

PASS HTTP 200
PASS status is ok
PASS mongodb is reachable


POST Step 2 — Admin Login
http://127.0.0.1:8000/api/v1/auth/login
200 - 860 ms - 1.488 KB 4

PASS HTTP 200
PASS access_token present
PASS refresh_token present
PASS token_type is bearer

GET Step 3 — Get My Profile (Admin)
http://127.0.0.1:8000/api/v1/users/me
200 - 73 ms - 452 B 4

PASS HTTP 200
PASS role is admin

Implementation done for Phase 2 - Ingestion & Chunking



IMPLEMENTATION_STATUS.md

```
1 # Implementation Status
2
3 Tracks phase-by-phase progress against
4 `Enterprise-AI-Knowledge-Platform-Blueprint.md` §23.
5 Each phase requires explicit approval (after
6 validation) before the next one starts.
7
8 | Phase | Scope | Status |
9 | --- | --- | --- |
10 | **1 - Foundation** | Auth/RBAC, tenant model,
11 | MongoDB schema, audit logging | Validated &
12 | Approved |
13 | **2 - Ingestion & Chunking** | File/web/Confluence
14 | ingestion, 8 chunking strategies, embeddings |
15 | Implemented - Awaiting Validation |
16 | 3 - Search & RAG | Semantic/keyword/hybrid search,
17 | RAG QA | Not started |
18 | 4 - Agentic & Comparison | LangGraph agentic
19 | retrieval, document comparison | Not started |
20 | 5 - Analytics & Governance | Dashboards, audit/
21 | governance, enforcement | Not started |
```


37/37 tests pass — all Phase 1 and Phase 2 tests green.

Phase 2 — Complete. Validation Steps

Here's how to validate Phase 2 against your live Atlas instance:


ctrl esc to attach selected text

python - backend



Validating the Phase 2 – Ingestion and Chunking via API

Uploaded the .txt file for ingestion



My Workspace

Search

Invites Upgrade

POST Step 1 — Admin Login Phase 2 - Ingestion POST Step 2 — Upload GET Step 3 — Poll Status GET New Request No environment

Phase 2 - Ingestion and Chunking Step 3 — Poll Status Until Embedded

GET http://127.0.0.1:8000/api/v1/documents/{{document_id}}/status

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

Hide auto-generated headers

Key	Value	Description	Bulk Edit	Presets	...
Cache-Control	no-cache				
Postman-Token	<calculated when request is sent>				
Host	<calculated when request is sent>				
User-Agent	PostmanRuntime/7.54.0				

Body Cookies Headers 4 Test Results

200 OK 247 ms 224 B Save Response

JSON Preview Visualize

```
1 {
2   "document_id": "6a33ae9d978b07836adb61226",
3   "status": "embedded",
4   "chunk_count": 2,
5   "error_message": null
6 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0 Import Complete

Globals Vault Tools

Type here to search

31°C 17:44 18-06-2026

Embeddings for the document had been validated

My Workspace

Search

Invites Upgrade

POST Step 1 — Admin Login Phase 2 - Ingestion POST Step 2 — Upload GET Step 3 — Poll Status GET Step 4 — List Documents No environment

Phase 2 - Ingestion and Chunking Step 4 — List Documents

GET http://127.0.0.1:8000/api/v1/documents

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

User-Agent PostmanRuntime/7.54.0

Accept */*

Accept-Encoding gzip, deflate, br

Connection keep-alive

Authorization Bearer {{token}}

Body Cookies Headers 4 Test Results

200 OK 315 ms 1.55 KB Save Response

JSON Preview Visualize

```
1 {
2   "tenant_id": "6a33a789eedf8f8c1184b71f",
3   "source_type": "txt",
4   "title": "Prompt for designing architecture.txt",
5   "original_filename": "Prompt for designing architecture.txt",
6   "storage_uri":
7     "C:\\Users\\hp\\Desktop\\RAGPROJECT\\backend\\uploads\\6a33a789eedf8f8c1184b71f\\e2299e933fe849b09f974f226370b3a6_Prompt for designing architecture.txt",
8   "owner_id": "6a33a78aeeef8f8c1184b720",
9   "classification": "internal",
10  "lineage": {
11    "ingested_by": "6a33a78aeeef8f8c1184b720",
12    "ingestion_job_id": "028d6c3c-ca4b-4eee-a984-a29dece1a2d2",
13    "source_url": null,
14    "source_system": null
15  }
16 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0 Import Complete

Globals Vault Tools

Type here to search

31°C 17:46 18-06-2026

My Workspace

Search

Invites Upgrade

POST Step 1 — Admin Login Phase 2 - Ingestion POST Step 2 — Upload GET Step 3 — Poll Status GET Step 4 — List Documents

Phase 2 - Ingestion and Chunking > Step 4 — List Documents

GET http://127.0.0.1:8000/api/v1/documents

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK • 315 ms • 1.55 KB Save Response

{ } JSON Preview Visualize

```
14 {
15   "source_system": null
16 },
17   "language": "en",
18   "status": "embedded",
19   "version": 1,
20   "is_deleted": false,
21   "chunk_count": 2,
22   "created_at": "2026-06-18T12:13:13.487000",
23   "updated_at": "2026-06-18T12:13:26.963000",
24   "id": "6a33e0d978b07836adb61226"
25 },
26 {
27   "tenant_id": "6a33a789eefd8f8c1184b71f",
28   "source_type": "txt",
29   "title": "Requirement prompt.txt",
30   "original_filename": "Requirement prompt.txt",
31   "storage_uri": "C:\\Users\\hp\\Desktop\\RAGPROJECT\\backend\\uploads\\6a33a789eefd8f8c1184b71f\\e2298e933fe849b09f974f226370b3a6_Prompt for designing architecture.txt",
32   "owner_id": "6a33a78aeefd8f8c1184b720",
33   "classification": "internal",
34   "lineage": {
35     "ingested_by": "6a33a78aeefd8f8c1184b720",
36     "ingestion_job_id": "028d6c3c-ca4b-4eee-a04a-a29dece1a2d2",
37     "source_uri": null,
38     "source_system": null
39   },
40   "language": "en",
41 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal Import Complete

Type here to search

31°C 17:47 18-06-2026

My Workspace

Search

Invites Upgrade

POST Step 1 — Admin Login Phase 2 - Ingestion POST Step 2 — Upload GET Step 3 — Poll Status GET Step 4 — List Documents GET New Request GET Get Single Document

Phase 2 - Ingestion and Chunking > Get Single Document

GET http://127.0.0.1:8000/api/v1/documents/{{document_id}}

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK • 524 ms • 876 B Save Response

{ } JSON Preview Visualize

```
1 {
2   "tenant_id": "6a33a789eefd8f8c1184b71f",
3   "source_type": "txt",
4   "title": "Prompt for designing architecture.txt",
5   "original_filename": "Prompt for designing architecture.txt",
6   "storage_uri": "C:\\Users\\hp\\Desktop\\RAGPROJECT\\backend\\uploads\\6a33a789eefd8f8c1184b71f\\e2298e933fe849b09f974f226370b3a6_Prompt for designing architecture.txt",
7   "owner_id": "6a33a78aeefd8f8c1184b720",
8   "classification": "internal",
9   "lineage": {
10     "ingested_by": "6a33a78aeefd8f8c1184b720",
11     "ingestion_job_id": "028d6c3c-ca4b-4eee-a04a-a29dece1a2d2",
12     "source_uri": null,
13     "source_system": null
14   },
15   "language": "en",
16 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal Import Complete

Type here to search

31°C 17:51 18-06-2026

Step 6 — Chunking Preview (dry-run, no DB write)

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 7 - Fetch Persisted Chunks

GET http://127.0.0.1:8000/api/v1/chunking/{{document_id}}/chunks

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results Save Response

200 OK 456 ms · 2.87 KB

```
{
  "tenant_id": "6a33a789eefd8f8c1184b71f",
  "document_id": "6a33e8d978b07836adb61226",
  "chunk_index": 0,
  "text": "ICE POT \r\n\r\nInstruction:\r\n\r\nHelp me to design architecture for testcase generator .use these input and output requirements for the same\r\n\r\nDefine API endpoints also\r\n\r\nInput:\r\n\r\nJira,userstory,Requirements,figma design\r\n\r\nConfluence,A00,browser stack\r\n\r\nWiki\r\n\r\nUX design,\r\n\r\nDefect Database\r\n\r\nGithub report\r\n\r\nProd tickets\r\n\r\nTechnical documentation\r\n\r\nInputOutput:\r\n\r\nTestcases (Specific format,type)\r\n\r\nCritical functionality\r\n\r\nInjestion Pipeline-injest all the existing testcases and userstories by creating embedding\r\n\r\nRetrieval pipeline:\r\n\r\nHybrid search\r\n\r\nRerank,Deduplication,\r\n\r\nSummarization\r\n\r\nPrompting\r\n\r\nLLM output\r\n\r\nVerify the output\r\n\r\nVector database :MongoDB \r\n\r\nembedding model:Mistral AI\r\n\r\nTechstack:\r\n\r\nFront end: react.js\r\n\r\nBackend :Node and typescript\r\n\r\nFeedback mechanism (Approve/reject/edit)\r\n\r\nFolder structure you can decide as per enterprise level and then design\r\n\r\nContext :\r\n\r\nYou are an chief architect in designing testcase generator\r\n\r\nPersona:"
```

Uploading PDF file for Ingestion

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 7 - Fetch Persisted Chunks

GET http://127.0.0.1:8000/api/v1/chunking/{{document_id}}/chunks

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results Save Response

200 OK 456 ms · 2.87 KB

```
{
  "tenant_id": "6a33a789eefd8f8c1184b71f",
  "document_id": "6a33e8d978b07836adb61226",
  "chunk_index": 1,
  "text": "folder structure you can decide as per enterprise level and then design\r\n\r\nContext :\r\n\r\nYou are an chief architect in designing testcase generator\r\n\r\nPersona:\r\n\r\nYou act an chief arctect to design a testcase generator\r\n\r\nOutput format:\r\n\r\n1.it should be in downloadable excel ,.csv,json,.etc format\r\n\r\n2.Should be in feature file format\r\n\r\n3.Tone :\r\n\r\nFormal and professional\r\n\r\n4.Start with MCQ questionnaire and ask me the relevant question to design the architecture document.\r\n\r\nnone question at a time and each question should have 4 options to choose.\r\n\r\nIncase if dont know the answer among these 4 ,i will let you know , then you can choose best one .\r\n\r\nIncase if i do not understand anything ,i will ask you to explain then after the explanation i will choose the answer.\r\n\r\n[STRICT]Based on my answers only you are going to design the architecture file and give me in the .md foformat\r\n\r\nDO NOT HALLUCINATE AND ASK ME THE PROPER QUESTIONS",
  "token_count": 201,
  "strategy_used": "recursive",
  "metadata": {

```

PDF Uploaded

My Workspace

Search

Phase 2 - Ingestion and Chunking > Step 8 — Upload a PDF

POST http://127.0.0.1:8000/api/v1/documents/upload

Body

form-data

Key	Value	Description	Bulk Edit
file	MongoDB-Explanation-for-QATest-Engineers (1).pdf		

Body

```
{
  "document_id": "6a33f91578b87836adb6122b",
  "status": "processing",
  "message": "Ingestion started."
}
```

202 Accepted · 1.18 s · 226 B

Activate Windows
Go to Settings to activate Windows.

document is uploaded for ingestion

My Workspace

Search

Phase 2 - Ingestion and Chunking > Step 9 — Upload a DOCX

POST http://127.0.0.1:8000/api/v1/documents/upload

Body

form-data

Key	Value	Description	Bulk Edit
file	Business Requirement.docx		


Body

```
{
  "document_id": "6a33f95578b87836adb61235",
  "status": "processing",
  "message": "Ingestion started."
}
```

202 Accepted · 469 ms · 226 B

Activate Windows
Go to Settings to activate Windows.

Uploaded the Unsupported file and checking for an error



Running different chunking categories

Strategy	Expected
fixed	200, chunks of roughly equal character length
recursive	200, chunks split at natural paragraph/sentence breaks
semantic	200, chunks grouped by semantic similarity
structure	200, chunks split by headings
query_aware	200, chunks around Q&A patterns
metadata_aware	200, table rows kept intact
llm_based	200, falls back to recursive (no GROQ_API_KEY set)
agentid	200, falls back to recursive (no GROQ_API_KEY set)

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Body

```
2 {
3   "document_id": "{{document_id}}",
4   "strategy": "fixed",
5   "chunk_size": 500,
6   "chunk_overlap": 50,
7   "overwrite": true
8 }
```

Body Cookies Headers 4 Test Results

202 Accepted 337 ms 214 B

Save Response

JSON

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "fixed",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 2 Import Complete

Globals Vault Tools

Type here to search

29°C

20:04
18-06-2026

2) Recursive

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Body

```
2 {
3   "document_id": "{{document_id}}",
4   "strategy": "recursive",
5   "chunk_size": 500,
6   "chunk_overlap": 50,
7   "overwrite": true
8 }
```

Body Cookies Headers 4 Test Results

202 Accepted 868 ms 218 B

Save Response

JSON

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "recursive",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 2 Import Complete


Globals Vault Tools

Type here to search


26°C

20:05
18-06-2026

3) semantic



4) Structure



5) query_aware

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Docs Params Authorization Headers 11 Body Scripts Settings

none form-data x-www-form-urlencoded raw binary GraphQL JSON

```
2 {
3   "document_id": "{{document_id}}",
4   "strategy": "query_aware",
5   "chunk_size": 500,
6   "chunk_overlap": 50,
7   "overwrite": true
8 }
```

Body Cookies Headers 4 Test Results

202 Accepted 184 ms 220 B

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "query_aware",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

26°C

20:08
18-06-2026

metadata_aware

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Docs Params Authorization Headers 11 Body Scripts Settings

none form-data x-www-form-urlencoded raw binary GraphQL JSON

```
2 {
3   "document_id": "{{document_id}}",
4   "strategy": "metadata_aware",
5   "chunk_size": 500,
6   "chunk_overlap": 50,
7   "overwrite": true
8 }
```

Body Cookies Headers 4 Test Results

202 Accepted 489 ms 223 B

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "metadata_aware",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

26°C

20:09
18-06-2026

LLM Based

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Body

```
1 {
2   "document_id": "{{document_id}}",
3   "strategy": "llm_based",
4   "chunk_size": 500,
5   "chunk_overlap": 50,
6   "overwrite": true
7 }
```

Body Cookies Headers 4 Test Results 202 Accepted 182 ms 218 B Save Response

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "llm_based",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

26°C

20:09
18-06-2026

Agentic

My Workspace

Search

Invite Upgrade

Phase 2 - Ingestion and Chunking > Step 12 — Re-chunk with Different Strategy

POST http://127.0.0.1:8000/api/v1/chunking/run

Body

```
1 {
2   "document_id": "{{document_id}}",
3   "strategy": "agentic",
4   "chunk_size": 500,
5   "chunk_overlap": 50,
6   "overwrite": true
7 }
```

Body Cookies Headers 4 Test Results 202 Accepted 301 ms 216 B Save Response

```
1 {
2   "document_id": "6a33e0d978b07836adb61226",
3   "strategy": "agentic",
4   "status": "processing"
5 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

26°C

20:10
18-06-2026

Deleted the existing document

The screenshot shows a REST client interface with a sidebar on the left listing various API steps. The main panel displays a DELETE request to the endpoint `http://127.0.0.1:8000/api/v1/documents/{{document_id}}`. The request body is empty, and the response is a 204 No Content status. The interface includes tabs for Docs, Params, Authorization, Headers, Body, Scripts, and Settings. The Body tab is active, showing an empty JSON object. The response tab shows a 204 No Content status with a response time of 284 ms and a body size of 113 B.

Phase 2 - Ingestion and Chunking > Step 14 — Soft Delete

DELETE `http://127.0.0.1:8000/api/v1/documents/{{document_id}}` Send

Docs Params Authorization Headers 8 Body Scripts Settings

☒ none ☐ form-data ☐ x-www-form-urlencoded ☐ raw ☐ binary ☐ GraphQL

This request does not have a body

Body Cookies Headers 3 Test Results 204 No Content · 284 ms · 113 B Save Response

{ } JSON Preview Visualize

1 |

Activate Windows
Go to Settings to activate Windows.

Proceeding to Phase 3

created vector search

The screenshot shows the MongoDB Atlas Search & Vector Search interface. The left sidebar contains navigation options like Search Indexes, AUTO EMBEDDING, Usage, Rate Limits, SHORTCUTS, Cluster Metrics, and Data Explorer. The main panel displays a table of search indexes. The table has columns for Database, Collection, Index Name, Status, Queryable, Type, Index Fields, Documents (Estimated), Size, and Required Memory. The table shows one index named `chunks_embedding_index` in the `chunks` collection of the `RAGTest` database. The index is in a `READY` status and is queryable. The index fields are `"embedding"`, `"tenant_id"`, `"document_id"`, and `"metadata.classification"`. The documents estimated are 0 (0%) indexed of 26. The size is 0B and the required memory is 0B.

cloud.mongodb.com/v2/6a3177f334ca16c3c2f34d81#/clusters/atlasSearch/RAG

Search & Vector Search

Find a search index

VIEW COMPOUND QUERY EXAMPLE CREATE SEARCH INDEX

10 free clusters have a limit of 3 search and vector indexes per cluster. Indexes Used: 1 of 3.

Database	Collection	Index Name	Status	Queryable	Type	Index Fields	Documents (Estimated)	Size	Required Memory
RAGTest	chunks	chunks_embedding_index	READY	✓	vectorSearch	"embedding", "tenant_id", "document_id", "metadata.classification"	0 (0%) indexed of 26	0B	0B

Learn more about Atlas Search
Learn more about Atlas Vector Search

Alert System Status: All Good
©2026 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales

Activate Windows
Go to Settings to activate Windows.

Started Implementing Phase 3 — Semantic Search & RAG

File Edit Selection View Go Run ... RAGPROJECT

EXPLORER

- ▼ RAGPROJECT
 - ▼ .claude
 - ▼ backend
 - ▼ .pytest_cache
 - ▼ app
 - ▼ keys
 - ▼ scripts
 - ▼ tests
 - ▼ uploads
 - ▼ venv
 - ▼ .env
 - ▼ .env.example
 - ▼ .gitignore
 - ▼ Dockerfile
 - ▼ pytest.ini
 - ▼ README.md
 - ▼ requirements.txt
 - ▼ Enterprise-AI-Knowled...
 - ▼ IMPLEMENTATION_ST...

IMPLEMENTATION_STATUS.md

```
1 # Implementation Status
2
3 Tracks phase-by-phase progress against
4 'Enterprise-AI-Knowledge-Platform-Blueprint.md' §23.
5 Each phase requires explicit approval (after
6 validation) before the next one starts.
7
8 | Phase | Scope | Status |
9 | --- | --- | --- |
10 | **1 - Foundation** | Auth/RBAC, tenant model,
11 MongoDB schema, audit logging | Validated &
12 Approved |
13 | **2 - Ingestion & Chunking** | File/web/Confluence
14 ingestion, 8 chunking strategies, embeddings |
15 Implemented - Awaiting Validation |
16 | 3 - Search & RAG | Semantic/keyword/hybrid search,
17 RAG QA | Not started |
18 | 4 - Agentic & Comparison | LangGraph agentic
19 retrieval, document comparison | Not started |
20 | 5 - Analytics & Governance | Dashboards, audit/
21 governance enforcement | Not started |
```

Implement Enterprise AI Knowledge Platfo...

image.png 1366x768

while clicking create search index button it says Please define the mapping document in your index structure

The button will be enabled this time because Vector Search uses the `fields` format, not `mappings`.

Proceed Phase 3 — Semantic Search & RAG

* Inferring...

Queue another message...

+ 1 line selected </> Edit automatically

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

(venv) PS C:\Users\hvp\Desktop\RAGPROJECT\backend> cd C:\Users\hvp\Desktop\RAGPROJECT\backend

>> .\venv\Scripts\Activate.ps1

>> python scripts/setup_vector_index.py

raise RuntimeError(f"Atlas API returned {resp.status_code}: {resp.text}")

RuntimeError: Atlas API returned 401: {

"error": 401,

"reason": "Unauthorized",

"detail": "You are not authorized for this resource."

>> (venv) PS C:\Users\hvp\Desktop\RAGPROJECT\backend>

Activate Windows

Go to Settings to activate Windows.

Chat quota reached

21:02 18-06-2026

Phase 3 implementation is done

File Edit Selection View Go Run ... RAGPROJECT

EXPLORER

- ▼ RAGPROJECT
 - ▼ .claude
 - ▼ backend
 - ▼ .pytest_cache
 - ▼ app
 - ▼ keys
 - ▼ scripts
 - ▼ tests
 - ▼ uploads
 - ▼ venv
 - ▼ .env
 - ▼ .env.example
 - ▼ .gitignore
 - ▼ Dockerfile
 - ▼ pytest.ini
 - ▼ README.md
 - ▼ requirements.txt
 - ▼ Enterprise-AI-Knowled...
 - ▼ IMPLEMENTATION_ST...

Enterprise-AI-Knowledge-Platform-Blueprint.md

```
1 # Enterprise AI Knowledge Platform - Implementation
2
3 ## 28. Learning Outcomes
4
5 - Structuring a **modular monolith** that preserves
6 a clean extraction path to microservices, matched to
7 actual (not assumed) concurrency needs.
8
9 - Using **LangGraph for bounded agentic workflows**
10 - gaining the benefits of multi-step reasoning while
11 explicitly avoiding open-ended autonomy that wasn't
12 requested.
13
14 ---
15
16 *End of Blueprint.*
```

Implement Enterprise AI Knowledge Platfo...

Proceed Phase 3 — Semantic Search & RAG

Phase 3 is complete. Here's what was built:

Phase 3 — What was implemented

Component	File	Details
Semantic Search	search/semantic.py	Atlas \$vectorSearch with tenant + classification filters

Proceed Phase 3 — Semantic Search & RAG

+ Enterprise-AI-Knowledge-Platform-Blu... </> Edit automatically

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

(venv) PS C:\Users\hvp\Desktop\RAGPROJECT\backend> cd C:\Users\hvp\Desktop\RAGPROJECT\backend

>> .\venv\Scripts\Activate.ps1

>> python scripts/setup_vector_index.py

raise RuntimeError(f"Atlas API returned {resp.status_code}: {resp.text}")

RuntimeError: Atlas API returned 401: {

"error": 401,

"reason": "Unauthorized",

"detail": "You are not authorized for this resource."

>> (venv) PS C:\Users\hvp\Desktop\RAGPROJECT\backend>

Activate Windows

Go to Settings to activate Windows.

Chat quota reached

21:32 18-06-2026

Created Atlas Search in MongoDB

The screenshot displays the MongoDB Atlas Search & Vector Search interface. The left sidebar shows navigation options: Search & Vector Search, Search Indexes, AUTO EMBEDDING, Usage, Rate Limits, SHORTCUTS, Cluster Metrics, and Data Explorer. The main content area is titled "Search & Vector Search" and includes a search bar and a "CREATE SEARCH INDEX" button. A notification states: "M0 free clusters have a limit of 3 search and vector indexes per cluster. Indexes Used: 2 of 3." Below this, a table lists the search indexes:

Database	Collection	Index Name	Status	Queryable	Type	Index Fields	Documents (Estimated)	Size	Required Memory
RAGTest	chunks	chunks_embedding_index	READY	✓	vectorSearch	"embedding", "tenant_id", "document_id", "metadata.classification"	26 (100%) indexed of 26	99.61KB	99.61KB
RAGTest	chunks	chunks_text_index	PENDING	✗	search	"metadata", "tenant_id", "text"	0 (0%) indexed of 26	0B	-

At the bottom of the interface, there is an "Alert System Status: All Good" message and an "Activate Windows" watermark.


This screenshot shows the same MongoDB Atlas Search & Vector Search interface, but with an update: the "chunks_text_index" is now in a "READY" state, and its "Queryable" status is marked with a green checkmark. The table data is as follows:

Database	Collection	Index Name	Status	Queryable	Type	Index Fields	Documents (Estimated)	Size	Required Memory
RAGTest	chunks	chunks_embedding_index	READY	✓	vectorSearch	"embedding", "tenant_id", "document_id", "metadata.classification"	26 (100%) indexed of 26	99.61KB	99.61KB
RAGTest	chunks	chunks_text_index	READY	✓	search	"metadata", "tenant_id", "text"	26 (100%) indexed of 26	0B	-


The rest of the interface, including the sidebar and bottom status bar, remains the same as in the previous screenshot.

Validating Phase 3

Restarted the server



Step 2 — Semantic Search



My Workspace

Search

Invite Upgrade

POST Step 1 — Login (get token) Phase 3 - Search & RAG GET Step 2 — Semantic Search

Phase 3 - Search & RAG > Step 2 — Semantic Search

GET http://127.0.0.1:8000/api/v1/search?q=machine+learning&mode=semantic&top_k=5

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

User-Agent PostmanRuntime/7.54.0

Accept */*

Accept-Encoding gzip, deflate, br

Connection keep-alive

Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 1 m 20.10 s 6.74 KB Save Response

JSON Preview Visualize

```
12 {
13   "score": 0.6542938351631165,
14   "strategy_used": "recursive",
15   "metadata": {
16     "classification": "internal",
17     "section_heading": null,
18     "page_number": null,
19     "source_url": null,
20     "block_type": null,
21     "score": null
22   }
23 },
24 {
25   "chunk_id": "6a33f96d78b07836adb6123a",
26   "document_id": "6a33f95578b07836adb61235",
27   "document_title": "Business Requirement.docx",
28   "chunk_index": 4,
29   "text": "10. Deployment Documentation\nSuccess Criteria:\n- Users can upload and process documents.
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Windows Taskbar: Type here to search, 28°C, 21:47, 18-06-2026

My Workspace

Search

Invite Upgrade

POST Step 1 — Login (get token) Phase 3 - Search & RAG GET Step 2 — Semantic Search

Phase 3 - Search & RAG > Step 2 — Semantic Search

GET http://127.0.0.1:8000/api/v1/search?q=machine+learning&mode=semantic&top_k=5

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

User-Agent PostmanRuntime/7.54.0

Accept */*

Accept-Encoding gzip, deflate, br

Connection keep-alive

Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 1 m 20.10 s 6.74 KB Save Response

JSON Preview Visualize

```
80 \nAnalytics Dashboard:\n- Total documents ingested.\n- Total chunks generated.\n- Total
81 embeddings created.",
82 "score": 0.6284177383314289,
83 "strategy_used": "recursive",
84 "metadata": {
85   "classification": "internal",
86   "section_heading": null,
87   "page_number": null,
88   "source_url": null,
89   "block_type": null,
90   "score": null
91 },
92 "total": 5,
93 "latency_ms": 78234.73
94 }
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Windows Taskbar: Type here to search, 28°C, 21:47, 18-06-2026

Step 3 : Keyword Search

Postman interface showing Step 3 - Keyword Search. The request is a GET to `http://127.0.0.1:8000/api/v1/search?q=machine+learning&mode=keyword&top_k=5&`. The response is a 200 OK with a JSON body:

```
{  "query_id": "4851b64a-6ecf-4fe9-80c6-6ee835c92c6f",  "query": "machine learning",  "mode": "keyword",  "results": [],  "total": 0,  "latency_ms": 99.66}
```

Postman interface showing Step 3 - Keyword Search. The request is a GET to `http://127.0.0.1:8000/api/v1/search?q=machine+learning&mode=keyword&top_k=5`. The response is a 200 OK with a JSON body:

```
{  "query_id": "499cff09-2e0c-4243-a121-5f0701a6cd16",  "query": "machine learning",  "mode": "keyword",  "results": [    {      "chunk_id": "6a341dde3bce5d64cd2235cc",      "document_id": "6a341dc53bce5d64cd2235cb",      "document_title": "ML.txt",      "chunk_index": 0,      "text": "Machine learning is a subset of artificial intelligence that enables computers to learn from data.\\nDeep learning uses neural networks with multiple layers to process complex patterns.\\nNatural language processing allows machines to understand and generate human language.\\nLarge language models like GPT and LLaMA are trained on massive text datasets.\\nVector embeddings represent text as numerical arrays for semantic similarity search.\\nRetrieval Augmented Generation combines search with language models to answer questions.\\nEnterprise knowledge management helps organizations store and retrieve institutional knowledge."    }  ]}
```

Hybrid Search

The screenshot shows the Postman interface with a workspace named "My Workspace". The left sidebar lists collections, with "Phase 3 - Search & RAG" expanded. The main panel displays a GET request to `http://127.0.0.1:8000/api/v1/search?q=artificial+intelligence&mode=hybrid&top_k=5`. The request is successful (200 OK) with a response time of 56.83 s and a body size of 6.48 KB. The response body is a JSON object containing a query ID, query, mode, and a list of results. The first result is a document titled "ML.txt" with a chunk index of 0. The text of the document is: "Machine learning is a subset of artificial intelligence that enables computers to learn from data. Deep learning uses neural networks with multiple layers to process complex patterns. Natural language processing allows machines to understand and generate human language. Large language models like GPT and LLaMA are trained on massive text datasets. Vector embeddings represent text as numerical arrays for semantic similarity search. Retrieval Augmented Generation combines search with language models to answer questions. Enterprise knowledge management helps organizations store and retrieve institutional knowledge."

```
{
  "query_id": "bb4ba664-c3bd-40e9-8530-5782e3d8317f",
  "query": "artificial intelligence",
  "mode": "hybrid",
  "results": [
    {
      "chunk_id": "6a341dde3bce5d64cd2235cc",
      "document_id": "6a341dc53bce5d64cd2235cb",
      "document_title": "ML.txt",
      "chunk_index": 0,
      "text": "Machine learning is a subset of artificial intelligence that enables computers to learn from data. Deep learning uses neural networks with multiple layers to process complex patterns. Natural language processing allows machines to understand and generate human language. Large language models like GPT and LLaMA are trained on massive text datasets. Vector embeddings represent text as numerical arrays for semantic similarity search. Retrieval Augmented Generation combines search with language models to answer questions. Enterprise knowledge management helps organizations store and retrieve institutional knowledge."
    }
  ]
}
```

The screenshot shows the Postman interface with a workspace named "My Workspace". The left sidebar lists collections, with "Phase 3 - Search & RAG" expanded. The main panel displays a GET request to `http://127.0.0.1:8000/api/v1/search?q=test&mode=hybrid`. The request is successful (200 OK) with a response time of 12.99 s and a body size of 11.4 KB. The response body is a JSON object containing a query ID, query, mode, and a list of results. The first result is a document titled "MongoDB-Explanation-for-QATest-Engineers (1).pdf" with a chunk index of 7. The text of the document is: "Why MongoDB is Powerful for QA & AI\nTest data is inherently dynamic and unpredictable. Different test cases contain different fields, different APIs return different structures, and different validations require different data shapes. MongoDB was built for exactly this reality. Flexible Schema\nNo two test cases need to look the same. One TC can have 3 fields, another can have 15. MongoDB handles both without any schema migration. Native JSON/API Support\nAPI responses are already JSON. MongoDB stores them directly (no transformation, no mapping, no friction between your test and your database). Native RAG System"

```
{
  "query_id": "13780a79-c793-4524-aabc-cdb186e54a5d",
  "query": "test",
  "mode": "hybrid",
  "results": [
    {
      "chunk_id": "6a33f93778b87836adb61233",
      "document_id": "6a33f91578b87836adb6122b",
      "document_title": "MongoDB-Explanation-for-QATest-Engineers (1).pdf",
      "chunk_index": 7,
      "text": "Why MongoDB is Powerful for QA & AI\nTest data is inherently dynamic and unpredictable. Different test cases contain different fields, different APIs return different structures, and different validations require different data shapes. MongoDB was built for exactly this reality. Flexible Schema\nNo two test cases need to look the same. One TC can have 3 fields, another can have 15. MongoDB handles both without any schema migration. Native JSON/API Support\nAPI responses are already JSON. MongoDB stores them directly (no transformation, no mapping, no friction between your test and your database). Native RAG System"
    }
  ]
}
```

My Workspace

Search

Invite Upgrade

Phase 3 - Search & RAG > New Request

GET http://127.0.0.1:8000/api/v1/search?q=test&mode=hybrid

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer ((alice_token))

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 12.99 s 11.4 KB Save Response

```
{
  "score": 0.014925373134328358,
  "strategy_used": "structure",
  "metadata": {
    "classification": "internal",
    "section_heading": null,
    "page_number": null,
    "source_url": null,
    "block_type": null,
    "score": null
  }
},
{
  "total": 10,
  "latency_ms": 12119.03
}
```

Activate Windows
Go to Settings to activate Windows.

Step 8 — RAG Ask (Question Answering)

My Workspace

Search

Invite Upgrade

Phase 3 - Search & RAG > Step 8 — RAG Ask (Question Answering)

POST http://127.0.0.1:8000/api/v1/rag/ask

Send

Docs Params Authorization Headers 11 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer ((alice_access_token))

✓ Content-Type application/json

Body Cookies Headers 4 Test Results

200 OK 8.56 s 3.1 KB Save Response

```
{
  "session_id": "f7cbfd9-3994-4fd4-a816-d6a9817f7811",
  "question": "What is this document about?",
  "answer": "This document appears to be about MongoDB, its components such as collections and documents [1], and MongoDB Atlas, a cloud-hosted version of MongoDB [2]. It also touches on the requirements for an Enterprise AI Knowledge Platform [3], [4], [5], including its expected outcome and success criteria.",
  "citations": [
    {
      "citation_index": 1,
      "chunk_id": "6a33f93778b07836adb61232",
      "document_id": "6a33f91578b07836adb6122b",
      "document_title": "MongoDB-Explanation-for-QATest-Engineers (1).pdf",
      "text_excerpt": "Collections & Documents\nCollections and documents are the heart of MongoDB. Understanding them is essential for every QA engineer working with test data.\nWhat is a Collection?\nA collection is a group of similar documents - equivalent to a table in SQL. Collections don't enforce a fixed schema, so each document can have a different structure."
    }
  ]
}
```

Activate Windows
Go to Settings to activate Windows.

My Workspace

Search

Phase 3 - Search & RAG > POST Step 1 — Login (get token) POST Step 8 — RAG Ask (Que GET Step 9 — Get RAG Sessio

Phase 3 - Search & RAG > Step 9 — Get RAG Session

GET http://127.0.0.1:8000/api/v1/rag/sessions/ {{session_id}}

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK · 1.55 s · 3.14 KB Save Response

JSON Preview Visualize

Add request

```
{
  "session_id": "f7cbfd9-3994-4fd4-a816-d6a9817f7011",
  "question": "What is this document about?",
  "answer": "This document appears to be about MongoDB, its components such as collections and documents [1], and MongoDB Atlas, a cloud-hosted version of MongoDB [2]. It also touches on the requirements for an Enterprise AI Knowledge Platform [3], [4], [5], including its expected outcome and success criteria.",
  "citations": [
    {
      "citation_index": 1,
      "chunk_id": "6a33f93778b07836adb61232",
      "document_id": "6a33f91578b07836adb6122b",
      "document_title": "MongoDB-Explanation-for-QATest-Engineers (1).pdf",
      "text_excerpt": "Collections & Documents\nCollections and documents are the heart of MongoDB. Understanding them is essential for every QA engineer working with test data.\nWhat is a Collection?\nA collection is a group of similar documents - equivalent to a table in SQL. Collections don't enforce a fixed schema, so ea"
    }
  ]
}
```

Step 9 — Get RAG Session

My Workspace

Search

Phase 3 - Search & RAG > POST Step 1 — Login (get token) POST Step 8 — RAG Ask (Que GET Step 9 — Get RAG Sessio

Phase 3 - Search & RAG > Step 9 — Get RAG Session

GET http://127.0.0.1:8000/api/v1/rag/sessions/ {{session_id}}

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK · 1.55 s · 3.14 KB Save Response

JSON Preview Visualize

Add request

```
{
  "document_title": "Business Requirement.docx",
  "text_excerpt": "10. Deployment Documentation\nSuccess Criteria:\n- Users can upload and process documents.\n- Knowledge sources are searchable.\n- Retrieval results are accurate.\n- AI-generated summaries are meaningful.\n- Chunking strategies are visualized.\n- End-to-end RAG workflow operates successfully.\n- Platform is"
},
  "model_used": "llama-3.3-70b-versatile",
  "token_usage": {
    "prompt_tokens": 0,
    "completion_tokens": 0,
    "total_tokens": 0
  },
  "latency_ms": 8174.378099798394,
  "feedback": null,
  "created_at": "2026-06-18T20:34:44.030000"
}
```

Step 13 — Verify Search Query Logged

Testleaf x (2) WhatsApp x RAG Search | Cloud: MongoDB x API Keys - GroqCloud x Ask Gemini

cloud.mongodb.com/v2/6a3177f334ca16c3c2f34d81#/clusters/atlasSearch/RAG

Search & Vector Search

Find a search index

10 free clusters have a limit of 3 search and vector indexes per cluster. Indexes Used: 2 of 3.

Database	Collection	Index Name	Status	Queryable	Type	Index Fields	Documents (Estimated)	Size	Required Memory
RAGTest	chunks	chunks_embedding_index	READY	✓	vectorSearch	"embedding", "tenant_id", "document_id", "metadata.classification"	30 (100%) indexed of 30	131.4KB	131.4KB
RAGTest	chunks	chunks_text_index	READY	✓	search	"text"	30 (100%) indexed of 30	63.58KB	-

Learn more about Atlas Search

Learn more about Atlas Vector Search

Alert System Status: All Good

©2026 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales

Activate Windows
Go to Settings to activate Windows.

Type here to search

Testleaf x (2) WhatsApp x RAG Data | Cloud: MongoDB x API Keys - GroqCloud x Ask Gemini

cloud.mongodb.com/v2/6a3177f334ca16c3c2f34d81#/explorer/6a3178e30522f44f441d534/RAGTest/search_queries/find

Data Explorer

My Queries

Data Modeling

CLUSTERS (1)

Search clusters

RAG

RAGTest

- TestDB
- agentic_workflows
- audit_logs
- chunks
- document_comparisons
- documents
- notifications
- rag_sessions
- refresh_tokens
- search_queries
- tenants
- users
- usertvns

search_queries

Documents 32 Aggregations Schema Indexes 2 Validation

Type a query: { field: 'value' } or Generate query

EXPLAIN Reset Find Options

ADD DATA BULK EXPORT CODE

25 1 - 25 of 32

```
{
  "_id": ObjectId("6a3419ef3bce5d64cd2235be"),
  "query_id": "10f6d751-4687-4ee8-a944-3f15c86c3281",
  "tenant_id": "6a33a789eef08f8c1184b71f",
  "user_id": "6a33a789eef08f8c1184b720",
  "query_text": "machine learning",
  "query_type": "semantic",
  "result_count": 5,
  "top_k": 5,
  "latency_ms": 78234.72788012526,
  "clicked_chunk_ids": Array (empty),
  "created_at": 2026-06-18T16:16:47.378+00:00
}
```

```
{
  "_id": ObjectId("6a341ac53bce5d64cd2235bf"),
  "query_id": "4851b64a-6ecf-4fe9-88c6-6ee035c92c6f",
  "tenant_id": "6a33a789eef08f8c1184b71f",
  "user_id": "6a33a789eef08f8c1184b720",
  "query_text": "machine learning",
  "query_type": "keyword",
  "result_count": 8,
  "top_k": 5
}
```


Alert System Status: All Good

©2026 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales


Activate Windows
Go to Settings to activate Windows.

Type here to search


Step 14 — Verify RAG Session Logged




Validation of Phase 3 had completed so started implementing Phase 4 — Agentic Retrieval & Document Comparison.



Phase 4 Implementation were done



Validating Phase 4 Implementation via API



Validating Step 1 in Phase 4

My Workspace

Phase 4 — Agentic Retrieval & Document Comparison > Step 2 — Agentic Run (Question Answering)

POST http://127.0.0.1:8000/api/v1/agent/run

Body

```
1 {
2   "question": "What are the key business requirements?",
3   "mode": "hybrid",
4   "top_k": 5,
5   "max_iterations": 3
6 }
```

Body Cookies Headers 4 Test Results

200 OK • 1 m 49.32 s • 3.32 KB

JSON Preview Visualize

```
34 ],
35   "retrieved_chunks": 5,
36   "iterations": 1,
37   "status": "completed",
38   "graph_state_trace": [
39     {
40       "node": "plan",
41       "iteration": 0,
42       "retrieved_count": null,
43       "sufficient": null,
44       "notes": "sub_queries=['What are the key business requirements?']"
45     },
46     {
47       "node": "retrieve",
48       "iteration": 1,
49       "retrieved_count": 5,
50       "sufficient": null,
```

Activate Windows
Go to Settings to activate Windows.

My Workspace

Phase 4 — Agentic Retrieval & Document Comparison > Step 2 — Agentic Run (Question Answering)

POST http://127.0.0.1:8000/api/v1/agent/run

Body

```
1 {
2   "question": "What are the key business requirements?",
3   "mode": "hybrid",
4   "top_k": 5,
5   "max_iterations": 3
6 }
```

Body Cookies Headers 4 Test Results

200 OK • 1 m 49.32 s • 3.32 KB

JSON Preview Visualize

```
54   "node": "evaluate",
55   "iteration": 1,
56   "retrieved_count": 5,
57   "sufficient": true,
58   "notes": "heuristic: chunks>=2 or max_iter reached"
59 },
60 {
61   "node": "synthesize",
62   "iteration": 1,
63   "retrieved_count": null,
64   "sufficient": null,
65   "notes": "LLM answer generated"
66 },
67 ],
68 "model_used": "llama-3.3-70b-versatile",
69 "token_usage": {
70   "prompt_tokens": 0,
```

Activate Windows
Go to Settings to activate Windows.

My Workspace

Phase 4 — Agentic Retrieval & Document Comparison > Step 2 — Agentic Run (Question Answering)

POST http://127.0.0.1:8000/api/v1/agentic/run

Body

```
1 {
2   "question": "What are the key business requirements?",
3   "mode": "hybrid",
4   "top_k": 5,
5   "max_iterations": 3
6 }
```

Body Cookies Headers 4 Test Results 200 OK • 1 m 49.32 s • 3.32 KB

JSON Preview Visualize

```
59 {
60   {
61     "node": "synthesize",
62     "iteration": 1,
63     "retrieved_count": null,
64     "sufficient": null,
65     "notes": "LLM answer generated"
66   },
67 },
68 "model_used": "llama-3.3-70b-versatile",
69 "token_usage": {
70   "prompt_tokens": 0,
71   "completion_tokens": 0,
72   "total_tokens": 0
73 },
74 "latency_ms": 17716.65810002014
75 }
```

Activate Windows
Go to Settings to activate Windows.

Step 3 — Get Workflow by ID

My Workspace

Phase 4 — Agentic Retrieval & Document Comparison > Step 3 — Get Workflow by ID

GET http://127.0.0.1:8000/api/v1/agentic/workflows/{{workflow_id}}

Body Cookies Headers 4 Test Results 200 OK • 175 ms • 3.26 KB

JSON Preview Visualize

```
1 {
2   "workflow_id": "5f4ca090-d6a4-4fa5-9aee-dd7376883379",
3   "question": "What are the key business requirements?",
4   "answer": "I could not find a sufficient answer in the available documents. However, I can start the discovery process to identify the key business requirements [2], [4]. As a Business Analyst, I will ask questions to gather the necessary information [3]. Let's start with the project type [5].\n\nQuestion:\nWhat type of application are you planning to build?\nA. Enterprise RAG Platform\nB. AI Chatbot\nC. Document Intelligence Platform\nD. Workflow Automation System\nE. Skip",
5   "citations": [
6     {
7       "citation_index": 2,
8       "chunk_id": "6a33d1ab78b07836adb6121e",
9       "document_id": "6a33d14c78b07836adb6121c",
10      "document_title": "Requirement prompt.txt",
11      "text_excerpt": "2. Ask for confirmation.\r\n\r\n3. After confirmation, generate the complete implementation blueprint.\r\n\r\nThe final blueprint should include:\r\n\r\nBusiness Objectives\r\n\r\nArchitecture\r\n\r\nTechnology Stack\r\n\r\nDatabase Design\r\n\r\nAPI"

```

My Workspace

Search

Invite Upgrade

POST Step 1 — Admin Login (fresh token) Phase 4 — Agentic Retrieval & Document Comparison GET Step 3 — Get Workflow by ID GET New Request GET New Request GET New Request No environment

Phase 4 — Agentic Retrieval & Document Comparison > Step 3 — Get Workflow by ID

GET http://127.0.0.1:8000/api/v1/agentic/workflows/ {{workflow_id}}

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent	PostmanRuntime/7.54.0	
✓ Accept	*/*	
✓ Accept-Encoding	gzip, deflate, br	
✓ Connection	keep-alive	
✓ Authorization	Bearer {{admin_token}}	
Key	Value	Description

Body Cookies Headers 4 Test Results 200 OK 175 ms 3.26 KB Save Response

```
{
  "retrieved_chunks": 5,
  "iterations": 1,
  "status": "completed",
  "graph_state_trace": [
    {
      "node": "plan",
      "iteration": 0,
      "notes": "sub_queries=['What are the key business requirements?']"
    },
    {
      "node": "retrieve",
      "iteration": 1,
      "retrieved_count": 5,
      "notes": "merged 5 unique chunks"
    },
    {
      "node": "evaluate",

```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Windows Taskbar: Type here to search, 27°C, 03:25 19-06-2026

Step 4 — Workflow Not Found → 404

My Workspace

Search

Invite Upgrade

POST Step 1 — Admin Login (fresh token) Phase 4 — Agentic Retrieval & Document Comparison GET Step 3 — Get Workflow by ID GET New Request GET New Request GET New Request GET Step 4 — Workflow Not Found → 404 No environment

Phase 4 — Agentic Retrieval & Document Comparison > Step 4 — Workflow Not Found → 404

GET http://127.0.0.1:8000/api/v1/agentic/workflows/nonexistent-id

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent	PostmanRuntime/7.54.0	
✓ Accept	*/*	
✓ Accept-Encoding	gzip, deflate, br	
✓ Connection	keep-alive	
✓ Authorization	Bearer {{admin_token}}	
Key	Value	Description

Body Cookies Headers 4 Test Results 404 Not Found 1.53 s 164 B Save Response

```
{
  "detail": "Workflow not found."
}
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Windows Taskbar: Type here to search, 27°C, 03:27 19-06-2026

Step 5 — Document Comparison

My Workspace

Search

Invite Upgrade

POST Step 1 - Phase 4 - GET Step 3 - GET New Rec GET New Rec GET New Rec GET Step 4 - GET Step 5 - No environment

Phase 4 - Agentic Retrieval & Document Comparison > Step 5 - Document Comparison

GET http://127.0.0.1:8000/api/v1/documents

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{admin_token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 2.12 s 5.03 KB Save Response

JSON Preview Visualize

```
1 [
2   {
3     "tenant_id": "6a33a789eefd8f8c1184b71f",
4     "source_type": "docx",
5     "title": "Business Requirement.docx",
6     "original_filename": "Business Requirement.docx",
7     "storage_uri":
8       "C:\\Users\\hp\\Desktop\\RAGPROJECT\\backend\\uploads\\6a33a789eefd8f8c1184b71f\\57c83dbb275344d7adba
9       f7267f630f08_Business Requirement.docx",
10    "owner_id": "6a33a78aeefd8f8c1184b720",
11    "classification": "internal",
12    "lineage": {
13      "ingested_by": "6a33a78aeefd8f8c1184b720",
14      "ingestion_job_id": "b5b525c1-fbda-4054-a86f-1ff90a2ff3eb",
15      "source_uri": null,
16      "source_system": null
17    }
18  }
19 ]
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

27°C

ENG 03:30 19-06-2026

My Workspace

Search

Invite Upgrade

POST Step 1 - Phase 4 - GET Step 3 - GET New Rec GET New Rec GET New Rec GET Step 4 - GET Step 5 - No environment

Phase 4 - Agentic Retrieval & Document Comparison > Step 5 - Document Comparison

GET http://127.0.0.1:8000/api/v1/documents

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

✓ User-Agent PostmanRuntime/7.54.0

✓ Accept */*

✓ Accept-Encoding gzip, deflate, br

✓ Connection keep-alive

✓ Authorization Bearer {{admin_token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 2.12 s 5.03 KB Save Response

JSON Preview Visualize

```
37   "source_system": null
38 },
39   "language": "en",
40   "status": "embedded",
41   "version": 1,
42   "is_deleted": false,
43   "chunk_count": 9,
44   "created_at": "2026-06-18T21:40:20.930000",
45   "updated_at": "2026-06-18T21:40:52.810000",
46   "id": "6a3465c49785d4e00ad41332"
47 },
48 {
49   "tenant_id": "6a33a789eefd8f8c1184b71f",
50   "source_type": "txt",
51   "title": "ML.txt",
52   "original_filename": "ML.txt",
53   "storage_uri":

```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 Import Complete

Globals Vault Tools

Type here to search

27°C

ENG 03:31 19-06-2026

My Workspace

Search

Invite Upgrade

POST Step 1- Phase 4 - GET Step 3- GET New Rec GET New Rec GET New Rec GET Step 4- GET Step 5- No environment

Phase 4 — Agentic Retrieval & Document Comparison > Step 5 — Document Comparison

GET http://127.0.0.1:8000/api/v1/documents

Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

User-Agent PostmanRuntime/7.54.0

Accept */*

Accept-Encoding gzip, deflate, br

Connection keep-alive

Authorization Bearer {{admin_token}}

Key Value Description

Body Cookies Headers 4 Test Results

200 OK 2.12 s 5.03 KB Save Response

```
{
  "classification": "internal",
  "lineage": {
    "ingested_by": "6a33a78aeef8c1184b720",
    "ingestion_job_id": "bfb31284-c3fe-4e4f-b440-7338c0320d89",
    "source_url": null,
    "source_system": null
  },
  "language": "en",
  "status": "embedded",
  "version": 1,
  "is_deleted": false,
  "chunk_count": 8,
  "created_at": "2026-06-18T11:06:52.183000",
  "updated_at": "2026-06-18T11:08:28.027000",
  "id": "6a33d14c78b07836adb6121c"
}
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0 Import Complete

Globals Vault Tools

Type here to search

27°C

ENG

03:31
19-06-2026

Step 5 b — Document Comparison

My Workspace

Search

Invite Upgrade

POST Step 1 — Admin L Phase 4 — Agentic POST Step 5 b — Docu GET New Request GET New Request + No environment

Phase 4 — Agentic Retrieval & Document Comparison > Step 5 b — Document Comparison

POST http://127.0.0.1:8000/api/v1/compare

Send

Docs Params Authorization Headers 11 Body Scripts Settings Cookies

none form-data x-www-form-urlencoded raw binary GraphQL JSON Schema Beautify

```
1 {
2   "document_id_a": "{{document_id_1}}",
3   "document_id_b": "{{document_id_2}}",
4 }
```

Body Cookies Headers 4 Test Results

200 OK 9.04 s 5.22 KB Save Response

```
1 {
2   "comparison_id": "d9872f4e-c29a-486b-bc41-a4ce0e91a997",
3   "document_id_a": "6a3466a29785d4e00ad4132d",
4   "document_id_b": "6a3466f29785d4e00ad4133c",
5   "title_a": "ML.txt",
6   "title_b": "Business Requirement.docx",
7   "overall_similarity": 0.6255,
8   "divergent_pairs": [
9     {
10      "chunk_id_a": "6a3466b99785d4e00ad4132e",
11      "chunk_id_b": "6a3466369785d4e00ad4133e",
12      "similarity": 0.4013,
13      "text_a": "Machine learning is a subset of artificial intelligence that enables computers to learn from data.\r\nDeep learning uses neural networks with multiple layers to process complex patterns.\r\nNatural language processing allows machines to understand and generate human language.\r\nLarge language models like",
14      "text_b": "6. Visualize chunking strategies for learning and analysis.\r\nEnable document
```

Activate Windows
Go to Settings to activate Windows.

Connect Git Console Terminal 0 0 0 Import Complete

Globals Vault Tools

Type here to search

27°C

ENG

03:35
19-06-2026

My Workspace

Search

POST Step 1 — Admin Login | Phase 4 — Agentic | POST Step 5 b — Document Comparison

Phase 4 — Agentic Retrieval & Document Comparison > Step 5 b — Document Comparison

POST http://127.0.0.1:8000/api/v1/compare

Body

```
1 {
2   "document_id_a": "{{document_id_1}}",
3   "document_id_b": "{{document_id_2}}"
4 }
```

Body

```
15 },
16 {
17   "chunk_id_a": "6a3465b99785d4e0ad4132f",
18   "chunk_id_b": "6a3466369785d4e0ad4133e",
19   "similarity": 0.4892,
20   "text_a": "Deep learning uses neural networks with multiple layers to process complex patterns.
21             \r\nNatural language processing allows machines to understand and generate human language.
22             \r\nLarge language models like GPT and LLaMA are trained on massive text datasets.
23             \r\nVector embeddings represent text as numerical ",
24   "text_b": "6. Visualize chunking strategies for learning and analysis.\n7. Enable document
25             comparison and similarity analysis.\n8. Support enterprise-scale retrieval and question
              answering.\n9. Provide an agent-based retrieval workflow for complex queries.\n10. Generate
              summaries and insights from retrieved conten"
```

200 OK • 9.04 s • 5.22 KB

Activate Windows
Go to Settings to activate Windows.

My Workspace

Search

POST Step 1 — Admin Login | Phase 4 — Agentic | POST Step 5 b — Document Comparison

Phase 4 — Agentic Retrieval & Document Comparison > Step 5 b — Document Comparison

POST http://127.0.0.1:8000/api/v1/compare

Body

```
1 {
2   "document_id_a": "{{document_id_1}}",
3   "document_id_b": "{{document_id_2}}"
4 }
```


Body

```
43 },
44 {
45   "llm_summary": "***Summary of key differences**:  
The two documents, ML.txt and Business Requirement.docx,  
have distinct focuses. Document A delves into technical aspects of machine learning, such as deep  
learning and natural language processing, whereas Document B appears to concentrate on business  
requirements and applications, including document comparison and retrieval. The overall similarity score  
of 0.63 suggests that while there may be some overlap, the documents have notable differences in content  
and emphasis.\n\n***Topics in Document A but not in Document B**:  
Document A covers specific machine  
learning concepts, including deep learning, neural networks, and natural language processing, which are  
not mentioned in Document B. These technical topics suggest that Document A is focused on the  
fundamentals of machine learning.\n\n***Topics in Document B but not in Document A**:  
Document B  
discusses business-oriented features, such as visualizing chunking strategies, document comparison,  
enterprise-scale retrieval, and agent-based retrieval workflows, which are not present in Document A.  
These topics imply that Document B is concerned with the practical applications and implementation of  
machine learning in a business context."
46   "latency_ms": 4748.681999980211
47 }
```


200 OK • 9.04 s • 5.22 KB

Activate Windows
Go to Settings to activate Windows.

Step 6 — Get Comparison by ID



Step 7 — Same Document Rejected → 400



Step 8 — Nonexistent Document → 404

Step 9 — Verify agentic_workflows in Atlas

Cloud MongoDB Explorer interface showing the 'agentio_workflows' collection. The left sidebar displays the 'Data Explorer' with a tree view of clusters and collections. The main panel shows a document with fields like 'workflow_id', 'tenant_id', 'user_id', 'question', 'answer', 'retrieved_chunk_ids', 'citations', 'iterations', 'status', 'graph_state_trace', and 'model_used'. The document content is displayed in a JSON format.

Document content (JSON):


```
{  "workflow_id": "5f4ca098-d6a4-4fa5-9aee-dd7376883379",  "tenant_id": "6a33a789eef8f8c1184b71f",  "user_id": "6a33a789eef8f8c1184b728",  "question": "What are the key business requirements?",  "answer": "I could not find a sufficient answer in the available documents. However...",  "retrieved_chunk_ids": Array (5),  "citations": Array (4),  "iterations": 1,  "status": "completed",  "graph_state_trace": Array (4),  "0": Object,    "node": "plan",    "iteration": 0,    "notes": "sub_queries=[\"What are the key business requirements?\"]",  "1": Object,    "node": "retrieve",    "iteration": 1,    "retrieved_count": 5,    "notes": "merged 5 unique chunks",  "2": Object,  "3": Object,  "model_used": "llama-3.3-70b-versatile"}
```

Step 10 — Verify document_comparisons in Atlas


Cloud MongoDB Explorer interface showing the 'document_comparisons' collection. The left sidebar displays the 'Data Explorer' with a tree view of clusters and collections. The main panel shows a document with fields like '_id', 'comparison_id', 'tenant_id', 'user_id', 'document_id_a', 'document_id_b', 'title_a', 'title_b', 'overall_similarity', 'divergent_pairs', and 'text_a'. The document content is displayed in a JSON format.

Document content (JSON):

```
{  "_id": ObjectId("6a346b89785d4e08ad41345"),  "comparison_id": "d9872f4e-c29a-486b-bc41-a4ce0e01a997",  "tenant_id": "6a33a789eef8f8c1184b71f",  "user_id": "6a33a789eef8f8c1184b728",  "document_id_a": "6a3465a29785d4e08ad4132d",  "document_id_b": "6a3465f29785d4e08ad4133c",  "title_a": "ML.txt",  "title_b": "Business Requirement.docx",  "overall_similarity": 0.6254703998565674,  "divergent_pairs": Array (5),  "0": Object,    "chunk_id_a": "6a3465b99785d4e08ad4132e",  "chunk_id_b": "6a3465369785d4e08ad4133e",  "similarity": 0.4813,  "text_a": "Machine learning is a subset of artificial intelligence that enables c...",  "text_b": "6. Visualize chunking strategies for learning and analysis.",  "7. Enable _",  "1": Object,    "chunk_id_a": "6a3465b99785d4e08ad4132f",  "chunk_id_b": "6a3465369785d4e08ad4133e",  "similarity": 0.4892,  "text_a": "Deep learning uses neural networks with multiple layers to process com..."}
```



Validation completed for Phase 4 so proceeding to implement Phase 5 - Analytics & Governance



Phase 5 implementation is done

My Workspace

Phase 5 - Analytics & Governance > Step 2 — Enterprise Dashboard (Admin)

GET http://127.0.0.1:8000/api/v1/analytics/dashboard

200 OK · 1 m 18.06 s · 1.93 KB

```
{
  "agentic_status_breakdown": [
    {
      "status": "completed",
      "count": 1
    }
  ],
  "comparison_total": 1,
  "classification_distribution": [
    {
      "classification": "internal",
      "count": 7
    }
  ],
  "recent_audit_logs": [
    {
      "action": "user.status_change",
      "timestamp": "2024-06-19T12:00:00Z",
      "user": "admin"
    }
  ]
}
```

Activate Windows
Go to Settings to activate Windows.

Step 3 — Retrieval Analytics (Admin)

My Workspace

Phase 5 - Analytics & Governance > Step 3 — Retrieval Analytics (Admin)

GET http://127.0.0.1:8000/api/v1/analytics/retrieval

200 OK · 366 ms · 405 B

```
{
  "click_through_rate": 0.8383,
  "dedup_rate": 0.0,
  "mode_breakdown": [
    {
      "mode": "hybrid",
      "count": 16
    },
    {
      "mode": "keyword",
      "count": 15
    },
    {
      "mode": "semantic",
      "count": 2
    }
  ]
}
```

Stabilize Step 2 500 error
Validate Step 3 response fields
Audit shared auth + base URL


Introducing Postman Cloud Agent
Run long-running tasks without babysitting your laptop. Start from Slack or from Postman. [Learn more](#)
Get started >

Describe what you need. Press @ for context, / for Skills.


Activate Windows
Go to Settings to activate Windows.

Step 4 — Employee Blocked from Dashboard → 403







Step 6 — GDPR Erasure: Create a Throwaway User




Step 7 — GDPR Erasure: Hard Delete




Step 8 — Verify User is Gone → 404




Step 9 — Employee Blocked from GDPR Erase → 403




Step 12 — Verify Ingestion Data in Dashboard



Validation for Phase 5 is completed and proceed to implement Phase 6 — Hardening & Launch.



Phase 6 is implemented



Step 1 — Health Endpoint Shows Phase 6 Version

POST Step 1 — Admin Login (fresh token) Copy

GET Step 2 — Enterprise Dashboard (Admin)

GET Step 3 — Retrieval Analytics (Admin)

POST Step 4 — Employee Blocked from Dashboard → 403

GET Step 4 — Employee Blocked from Dashboard → 403 b

GET Step 5 — Employee Blocked from Retrieval Analytics → 403

POST Step 6 — GDPR Erasure: Create a Throwaway User - 1

POST Step 7 — GDPR Erasure: Hard Delete

GET Step 8 — Verify User is Gone → 404

POST Step 9 — Employee Blocked from GDPR Erase → 403

POST Step 10 — GDPR Erase Nonexistent User → 404

GET Step 12 — Verify Ingestion Data in Dashboard

Phase 6 — Hardening & Launch.

POST Step 1 — Admin Login (fresh token)

POST Step 1 — Health Endpoint Shows Phase 6 Version

QA Bot TS - LangChain API Copy 2

RAG MongoDB - Data Ingestion

Resume Ingestion

ENVIRONMENTS

SPECS

FLAWS

Connect Git Console Terminal 0 0 0

Search

Invite Upgrade

Phase 6 — Hardening & Launch.

Step 1 — Health Endpoint Shows Phase 6 Version

GET http://127.0.0.1:8000/health

Send

Params Authorization Headers 7 Body Scripts Settings Cookies

Query Params

Key	Value	Description	Bulk Edit
Key	Value	Description	

Body JSON Preview Visualize

```
1 {
2   "status": "ok",
3   "mongodb": true,
4   "environment": "dev",
5   "version": "1.0.0-phase6"
6 }
```

200 OK • 121 ms • 534 B • Save Response

Run the full Phase 6 sequence

Extend auth reuse to other steps

Diagnose the 403/404 hardening steps

Introducing Postman Cloud Agent

Run long-running tasks without babysitting your laptop. Start from Slack or from Postman. [Learn more](#)

Get started

Describe what you need. Press @ for context, / for Skills.

Activate Windows

Go to Settings to activate Windows.

Globals Vault Tools 18:34 19-06-2026

Step 2 — Security Headers on Every Response

POST Step 1 — Admin Login (fresh token) Copy

GET Step 2 — Enterprise Dashboard (Admin)

GET Step 3 — Retrieval Analytics (Admin)

POST Step 4 — Employee Blocked from Dashboard → 403

GET Step 4 — Employee Blocked from Dashboard → 403 b

GET Step 5 — Employee Blocked from Retrieval Analytics → 403

POST Step 6 — GDPR Erasure: Create a Throwaway User - 1

POST Step 7 — GDPR Erasure: Hard Delete

GET Step 8 — Verify User is Gone → 404

POST Step 9 — Employee Blocked from GDPR Erase → 403

POST Step 10 — GDPR Erase Nonexistent User → 404

GET Step 12 — Verify Ingestion Data in Dashboard

Phase 6 — Hardening & Launch.

POST Step 1 — Admin Login (fresh token)

POST Step 1 — Health Endpoint Shows Phase 6 Version

GET Step 2 — Security Headers on Every Response

QA Bot TS - LangChain API Copy 2

RAG MongoDB - Data Ingestion

ENVIRONMENTS

SPECS

FLAWS

Connect Git Console Terminal 0 0 0

Search

Invite Upgrade

Phase 6 — Hardening & Launch.

Step 2 — Security Headers on Every Response

GET http://127.0.0.1:8000/health

Send

Params Authorization Headers 7 Body Scripts Settings Cookies

Headers

Key	Value
date	Fri, 19 Jun 2026 13:11:09 GMT
server	uvicorn
content-length	75
content-type	application/json
x-content-type-options	nosniff
x-frame-options	DENY
strict-transport-security	1; mode=block
referrer-policy	strict-origin-when-cross-origin
content-security-policy	max-age=63072000; includeSubDomains
permissions-policy	default-src 'none'; frame-ancestors 'none'

200 OK • 390 ms • 534 B • Save Response

Visualize Response as table

```
const response = pm.response.json()
const template = `
<table border="1">
  <thead>
    <tr>
      <th>Key</th>
      <th>Value</th>
    </tr>
  </thead>
  <tbody>
    <tr>
      <td>date</td>
      <td>Fri, 19 Jun 2026 13:11:09 GMT</td>
    </tr>
    <tr>
      <td>server</td>
      <td>uvicorn</td>
    </tr>
    <tr>
      <td>content-length</td>
      <td>75</td>
    </tr>
    <tr>
      <td>content-type</td>
      <td>application/json</td>
    </tr>
    <tr>
      <td>x-content-type-options</td>
      <td>nosniff</td>
    </tr>
    <tr>
      <td>x-frame-options</td>
      <td>DENY</td>
    </tr>
    <tr>
      <td>strict-transport-security</td>
      <td>1; mode=block</td>
    </tr>
    <tr>
      <td>referrer-policy</td>
      <td>strict-origin-when-cross-origin</td>
    </tr>
    <tr>
      <td>content-security-policy</td>
      <td>max-age=63072000; includeSubDomains</td>
    </tr>
    <tr>
      <td>permissions-policy</td>
      <td>default-src 'none'; frame-ancestors 'none'</td>
    </tr>
  </tbody>
</table>
`
pm.html(response, template)
`
```

Waiting for approval...

Introducing Postman Cloud Agent

Run long-running tasks without babysitting your laptop. Start from Slack or from Postman. [Learn more](#)


Get started

Describe what you need. Press @ for context, / for Skills.

Activate Windows

Go to Settings to activate Windows.

Globals Vault Tools 18:43 19-06-2026



Validation for Phase 6

